

# Machine Learning Interview Question and Answer SET-1

Submitted by:- Ambarish Singh

**1) What does one mean by the term “machine learning”?**

**Ans.1)**

Machine learning is a subfield of artificial intelligence that focuses on developing algorithms that enable computers to learn from data and improve their predictions over time, without being explicitly programmed.

**2) Can you think of 4 distinct types of issues where it shines?**

**Ans.2)**

- a) **Predictive modeling:** Predictive models can be built using machine learning algorithms to predict future events based on past data.
- b) **Image and speech recognition:** ML algorithms are capable of analyzing and recognizing patterns in visual and auditory data, enabling applications like image and speech recognition.
- c) **Natural Language Processing:** ML is used to perform language-related tasks such as sentiment analysis, translation, and text classification.

- d) **Anomaly detection:** Machine learning algorithms can identify anomalies and outliers in large datasets, making it useful for detecting fraud, intrusion, or other unusual behaviors.

### 3) What is a labeled training set, and how does it work?

#### Ans.3)

A labeled training set is a dataset used to train machine learning algorithms. The dataset contains input examples and their corresponding output labels, which serve as "answers" to the model. The algorithm uses the input/output pairs in the labeled training set to "learn" the relationships between inputs and outputs, and then make predictions on new, unseen inputs.

The quality of the training set is crucial to the performance of the machine learning model. A high-quality labeled training set will contain a diverse set of examples that accurately represent the problem that the model is trying to solve. The model can use these examples to generalize to new inputs and make accurate predictions.

### 4) What are the two most important tasks that are supervised?

#### Ans.4)

- a) **Classification:** This is a supervised task where the algorithm is trained to assign a label to a given input. For example, image classification, where the goal is to categorize an image into one of several predefined classes (e.g. cat, dog, flower).

- b) **Regression:** This is a supervised task where the goal is to predict a continuous value, such as a price, given a set of input features. For example, predicting the price of a house based on its size, location, and age.

## 5) Can you think of four examples of unsupervised tasks?

Ans.5)

- a) **Clustering:** This is an unsupervised task where the goal is to group similar instances together. For example, grouping customers into segments based on their purchasing habits.
- b) **Dimensionality reduction:** This is the task of reducing the number of features in a dataset while retaining as much information as possible. For example, reducing the number of pixels in an image while retaining its important features.
- c) **Anomaly detection:** This is the task of identifying instances in a dataset that deviate significantly from the norm. For example, detecting credit card fraud by identifying transactions that deviate from a customer's normal spending patterns.
- d) **Association rule mining:** This is the task of discovering relationships between variables in a dataset. For example, discovering associations between items purchased in a grocery store to identify frequently purchased items together.

**6) State the machine learning model that would be best to make a robot walk through various unfamiliar terrains?**

**Ans.6)**

**Reinforcement learning** would be the best machine learning model for making a robot walk through various unfamiliar terrains.

Reinforcement learning is a type of machine learning where an agent learns to take actions in an environment to maximize a reward signal. In the case of a walking robot, the agent (robot) would receive a reward signal for successfully navigating through the terrain, and the model would learn over time to choose actions that result in higher rewards. The model can then be applied to new, unseen terrains and the robot can learn to walk through them effectively.

**7) Which algorithm will you use to divide your customers into different groups?**

**Ans.7)**

**Clustering** is the type of machine learning algorithm that would be used to divide customers into different groups. Clustering algorithms group similar instances together and can be used for market segmentation, customer segmentation, and other use cases where the goal is to identify natural groupings within a dataset.

**K-Means** is a popular clustering algorithm that partitions data into a specified number of clusters. Another commonly used algorithm is Hierarchical Clustering, which builds a hierarchy of clusters by successively merging or splitting existing clusters. The choice of

algorithm will depend on the specific requirements and characteristics of the data.

**8) Will you consider the problem of spam detection to be a supervised or unsupervised learning Problem?**

**Ans.8)**

The problem of spam detection is typically considered a supervised learning problem. In supervised learning, the algorithm is trained on a labeled dataset where the input instances are labeled as either spam or not spam. The model then uses this information to learn the characteristics of spam messages and make predictions on new, unseen messages.

In the case of spam detection, the labeled training set would consist of email messages that have been manually labeled as spam or not spam. The model would use this information to learn the patterns and features that are indicative of spam messages and then use these patterns to make predictions on new, incoming messages.

**9) What is the concept of an online learning system?**

**Ans.9)**

Online learning is a type of machine learning where the algorithm receives data instances one at a time and updates its model after each instance, as opposed to batch learning, where the algorithm receives all instances at once. Online learning is useful for systems

that receive a continuous stream of data, such as web applications and recommendation systems.

In an online learning system, the algorithm uses the information from each new instance to update its model, rather than starting from scratch each time. This allows the algorithm to learn and adapt to changes in the data over time, without having to be retrained on the entire dataset. The result is a model that can evolve and improve over time as it processes more data.

### **10) What is out-of-core learning, and how does it differ from core learning?**

#### **Ans.10)**

Out-of-core learning is a type of machine learning where the algorithm can process datasets that are too large to fit into memory. It differs from in-core or "core" learning, where the algorithm loads the entire dataset into memory and processes it in a single batch.

In out-of-core learning, the algorithm processes the data in small chunks, loading only a portion of the data into memory at a time. This allows the algorithm to process datasets that are much larger than the available memory. Out-of-core learning is useful for datasets that are too large to fit into memory, such as big data and streaming data.

The main difference between out-of-core and in-core learning is that out-of-core learning can handle larger datasets, while in-core learning requires the entire dataset to fit into memory. Out-of-core learning is also slower than in-core learning, as the algorithm must read the data from disk, but it is more flexible and can handle much larger datasets.

**11) What kind of learning algorithm makes predictions using a similarity measure?**

**Ans.11)**

The kind of learning algorithm that makes predictions using a similarity measure is called instance-based learning or nearest neighbor learning.

Instance-based learning works by storing the training instances and comparing new instances to the stored instances using a similarity measure. The algorithm then makes a prediction based on the stored instance(s) that is most similar to the new instance. This can be done for regression and classification problems.

For example, in a classification problem, the nearest neighbor algorithm would find the stored instance that is most similar to a new instance, and predict the class of the new instance based on the class of the most similar stored instance. The similarity measure can be based on Euclidean distance, Manhattan distance, cosine similarity, or other metrics.

Instance-based learning is a simple and intuitive approach, but it can be computationally expensive and may not work well with high-dimensional data. Nevertheless, it can be a useful approach for small datasets or for problems where a human expert can easily identify the nearest neighbors.

**12) What's the difference between a model parameter and a hyperparameter in a learning algorithm?**

**Ans.12)**

In a machine learning algorithm, model parameters and hyperparameters are different types of values that control the behavior and performance of the model.

Model parameters are values that are learned from the training data during the training process. They define the structure and behavior of the model, and are used to make predictions on new, unseen data. Examples of model parameters include the weights in a neural network, the coefficients in a linear regression model, and the mean and standard deviation in a Gaussian distribution.

Hyperparameters, on the other hand, are values that are set prior to training and control aspects of the training process, such as the learning rate, the number of trees in a random forest, or the number of clusters in a k-means algorithm. Hyperparameters are not learned from the data, but are chosen by the practitioner based on their experience and the characteristics of the problem.

The difference between model parameters and hyperparameters is that model parameters are learned from the data, while hyperparameters are set by the practitioner and control the training process. The choice of hyperparameters can greatly affect the performance of the model, so it is important to choose them carefully. This is typically done through a process called hyperparameter tuning, where the practitioner tries different combinations of hyperparameters and evaluates the performance of the model on a validation set.



**13) What are the criteria that model-based learning algorithms look for ? What is the most popular method they use to achieve success? What method do they use to make predictions?**

**Ans.13)**

Model-based learning algorithms are a type of machine learning that build a mathematical model to represent the relationships between the features and target variable in the data. The criteria that model-based learning algorithms look for are:

Fit to the data: The model should fit the training data well, accurately capturing the relationships between the features and target variable.

Simplicity: The model should be simple and interpretable, with a small number of parameters that are easy to understand.

Generalization: The model should generalize well to new, unseen data, making accurate predictions on data it has not seen before.

The most popular method used by model-based learning algorithms to achieve success is regularization, which is a technique that penalizes complexity in the model. This helps prevent overfitting, which occurs when the model fits the training data too well and fails to generalize to new data.

For predictions, model-based learning algorithms typically use the mathematical model to make predictions on new, unseen data. The algorithm inputs the features of the new instance into the model, and the model returns a prediction for the target variable. This prediction can be used for tasks such as classification or regression.

Examples of model-based learning algorithms include linear regression, logistic regression, decision trees, and support vector machines.

**14) Can you name four of the most important Machine Learning challenges?**

**Ans.14)**

Yes, four of the most important challenges in machine learning are:

- a) Overfitting: This occurs when the model fits the training data too well, memorizing the training examples instead of learning the underlying patterns. Overfitting can result in poor performance on new, unseen data.
- b) Underfitting: This occurs when the model is too simple and does not capture the complexity of the data. This can result in poor performance on both the training and test data.
- c) Bias-Variance trade-off: This is the balance between having a model that is too complex (high variance) and a model that is too simple (high bias). Finding the right balance is important for good performance on new, unseen data.
- d) Data quality: The quality of the data can greatly affect the performance of the machine learning model. Issues such as missing values, outliers, and noisy data can impact the model's ability to learn the underlying patterns in the data.

These challenges can be addressed through techniques such as cross-validation, regularization, and feature engineering, but they remain a significant part of the machine learning process. Solving these challenges requires a good understanding of the data, the problem, and the model, as well as a strong experimentation and evaluation process.

**15) What happens if the model performs well on the training data but fails to generalize the results to new situations? Can you think of three different options?**

**Ans.15)**

If a machine learning model performs well on the training data but fails to generalize to new situations, it is said to be overfitting. Overfitting is a common problem in machine learning and can have significant consequences.

There are several options for addressing overfitting:

- a) **Regularization:** Regularization is a technique that penalizes complexity in the model, helping to prevent overfitting. This can be achieved through techniques such as L1 or L2 regularization.
- b) **Cross-Validation:** Cross-validation is a technique for evaluating the performance of a model by splitting the data into multiple folds and using each fold for testing and validation. This can help to identify overfitting and allow for more robust model selection.
- c) **Ensemble Methods:** Ensemble methods involve combining multiple models to make predictions, reducing the risk of

overfitting. This can be achieved through techniques such as bagging and boosting.

Each of these options has its own advantages and disadvantages, and the best approach will depend on the specific problem and data. It is important to carefully evaluate the performance of the model on both the training data and new, unseen data to avoid overfitting and ensure that the model generalizes well to new situations.

### **16) What exactly is a test set, and why would you need one?**

#### **Ans.16)**

A test set is a portion of the data that is set aside for evaluating the performance of a machine learning model. The purpose of a test set is to provide an unbiased evaluation of the model's performance on new, unseen data. The test set is used to simulate the real-world scenario of making predictions on new data.

The test set is used to determine the model's generalization performance, which is the ability of the model to make accurate predictions on new data that it has not seen before. The performance of the model on the test set provides an estimate of its performance on future, unseen data, which is a key consideration in many machine learning problems.

In order to ensure that the test set provides an accurate evaluation of the model's performance, it is important to use a different set of data than the data used for training the model. This can be achieved through techniques such as cross-validation or splitting the data into training, validation, and test sets.

Having a test set is a crucial part of the machine learning process, as it provides an objective measure of the model's performance and allows for the selection of the best model for the specific problem and data.

### **17) What is a validation set's purpose?**

#### **Ans.17)**

A validation set is a portion of the data that is set aside for evaluating the performance of a machine learning model during training. The purpose of a validation set is to provide an estimate of the model's performance on new, unseen data, and to tune the hyperparameters of the model.

The validation set is used to determine the generalization performance of the model, which is the ability of the model to make accurate predictions on new data that it has not seen before. The performance of the model on the validation set provides an estimate of its performance on future, unseen data, which is a key consideration in many machine learning problems.

In order to ensure that the validation set provides an accurate evaluation of the model's performance, it is important to use a different set of data than the data used for training the model. The validation set is used to adjust the hyperparameters of the model, such as the learning rate or the number of hidden units, to ensure that the model is optimized for the specific problem and data.

Having a validation set is a crucial part of the machine learning process, as it provides an estimate of the model's performance on new data and allows for the selection of the best model for the specific problem and data. The validation set is also used to tune the hyperparameters of the model, which can greatly impact its performance.

**18) What precisely is the train-dev kit, when will you need it, how do you put it to use?**

**Ans.18)**

The train-dev (development) set is a portion of the data that is set aside for evaluating the performance of a machine learning model during training. The train-dev set is similar to a validation set, but it is typically used in larger machine learning projects or when more computational resources are available.

The purpose of the train-dev set is to provide an estimate of the model's performance on new, unseen data, and to tune the hyperparameters of the model. The train-dev set is used to determine the generalization performance of the model, which is the ability of the model to make accurate predictions on new data that it has not seen before.

You would need a train-dev set when you have a large amount of data, and you want to use a portion of it for model selection and hyperparameter tuning, and another portion for final testing. This is useful in situations where the performance of the model on the validation set is not a good indicator of its performance on the test set.

To use the train-dev set, you would first split the data into three parts: training, validation, and testing. The training set is used to train the model, the validation set is used to adjust the hyperparameters, and the train-dev set is used to evaluate the model's performance and make decisions about model selection and hyperparameter tuning.

The train-dev set can be used to compare different models and select the best one, and to fine-tune the hyperparameters of the selected model. Once the model has been selected and the hyperparameters have been tuned, the model can be trained on the combined training and validation data and evaluated on the test set to obtain a final estimate of its performance on new, unseen data.

**19) What could go wrong if you use the test set to tune hyperparameters?**

**Ans.19)**

Using the test set to tune hyperparameters can result in overfitting, which is a type of modeling error that occurs when a model is too closely fit to the training data and does not generalize well to new, unseen data.

The test set is meant to be a representative sample of new, unseen data that the model will encounter in practice. If the test set is used to tune the hyperparameters of the model, then the model will be biased towards the test set, and its performance on the test set will not reflect its true generalization performance on new, unseen data.

Additionally, using the test set to tune the hyperparameters can lead to an over-optimistic estimate of the model's performance, since the test

set performance will likely be better than the performance on new, unseen data. This can result in a false sense of security and a lack of confidence in the model when it is deployed in the real world.

To avoid these problems, it is important to keep the test set separate and use a validation set or a train-dev set for hyperparameter tuning. This will ensure that the model is not overfitted to the test set, and that the performance of the model on the test set reflects its true generalization performance on new, unseen data.

-----THANK YOU -----